

Coping with Security / Safety Tensions in Low-End Embedded Devices

Gene Tsudik
CS Dept., UC Irvine

Joint work with: X. Carpent¹, K. Eldefrawy², N. Rattanavipanon¹, A. Sadeghi³

1 - UC Irvine, 2 - SRI International, 3 - TU Darmstadt

1

Roadmap

- Overview of Remote Attestation (RA)
- Problem Statement
 - Conflict between security of RA and real-time operation
- Tentative mitigation measures
 - Periodic self-measurements
 - Interruptible RA with shuffled measurements
 - Interruptible RA with memory locking
- Conclusions & future work

2

Internet-of-Things (IoT) Gadgets



3

IoT-specific Attacks On:

- **Sensing:** Privacy
- **Actuation:** Security & Safety
- **Either:** DDoS Sourcing (aka Zombification)

4

Constraints for Simple IoT Devices: large scale + low price

CPU Power	Battery
<p>Hard to prevent malware from entry Next best thing is to detect malware!</p>	
Memory Size	Hardware Cost

Plus, limited physical “real estate” and typically low security budget


5

Remote Attestation (RA)

- A means of detecting malware presence on an untrusted remote device
 - A security service
 - A protocol
 - An architecture
- All of the above...
- RA typically realized as an interaction between:
 - **Verifier:** trusted entity
 - **Prover:** untrusted (potentially infected) remote device
 - Goal: learn current internal (sw/fw) state of prover

6

Typical (on-demand) RA



RA Adversarial Model (DAC'16)


- Remote adversary
 - Exploits vulnerabilities to inject malware
- Local adversary
 - Controls communication channel(s)
- Physical adversary
 - Non-Intrusive: side-channel attacks
 - Stealthy Intrusive: “read-only”
 - Intrusive: alters hardware

8

Landscape of RA Techniques

- Purely Hardware-based
 - Dedicated hardware support, e.g., TPM, TrustZone, SGX
 - Overkill for simple or low-end IoT devices (cost and features)
- Purely Software-based
 - Relies on precise timing measurements, e.g., SCUBA, VIPER, PIONEER
 - Unrealistic assumptions for IoT except peripheral/legacy devices
- Hybrid
 - SW/HW co-design, e.g., SMART, TrustLite
 - Minimal hardware impact
 - **Seems like good fit for low-end devices**

9



10

Hybrid RA Techniques

- SMART [NDSS'12]+[DATE'14]
 - First hybrid design of RA for low-end microcontrollers
- TrustLite [EuroSys'14]
 - Supports secure interrupts
- TyTan [DAC'15]
 - TrustLite with real-time functionality
 - Process being measured cannot interrupt
- HYDRA [WiSec'17]
 - SMART implementation for medium-end devices (secure boot needed)
 - Formally verified *seL4* microkernel guarantees security properties
 - Especially, isolation
- VRASED ['18]
 - First formally verified RA design
 - Based on a version of SMART

11

RA **vs** Safety-Critical Operation

on SMART-based MSP430 @ 8MHz

RA takes:

≈4.5 seconds

to measure 48KB of flash

on HYDRA-based ODROID @ 2GHz

RA takes:

≈7 seconds

to measure 1GB of RAM

RA can interfere with safety-critical operation!

So... can we make RA interruptible?

12

Whither Interruptible RA Execution?

- Susceptibility to transient malware
 - Interrupts and erases itself **during** attestation
 - Avoids detection, leaves no trace
- Self-relocating malware
 - Interrupts and moves itself around **during** attestation
 - Avoids detection, **remains** on Prover
- Temporal inconsistency?
 - Memory can change during attestation
 - Computed measurement (e.g., MAC, HASH-to-be-signed) might reflect memory that never existed
 - Could be caused by malware or even benign software
 - An important issue **beyond** the RA context


13

TyTan: First Attempt

- Dynamically configurable execution-aware memory protection unit (EA-MPU) enforces:
 - Access rules to prover's attestation key
 - Immutability
- Real-time OS
 - Process can issue interrupt if it's not being attested
 - Provides isolation between processes
- Problems:
 - What if Verifier wants to attest a safety-critical process?
 - OS might be buggy → OS compromise can violate isolation → malware can move around during attestation

14

Solution Landscape



15


Memory Locking [AsiaCCS'18]

- Can we allow RA measurement process to be interruptible?
- While enforcing ***temporal consistency*** of measured memory
 - E.g., by locking memory segments = temporarily make them read-only
- Measurement (on prover) starts at time: t_S and ends at time: t_E

16


Memory Locking: All-Lock

- All memory is locked throughout attestation
- Consistency **from t_s to t_E**




Memory Locking: Dec-Lock

- All memory locked at the start
- A memory block is unlocked after it is measured
- Consistency **at t_s**



Memory Locking: Inc-Lock

- All memory unlocked at the start
- A memory block is locked after it is measured
- Consistency **at t_E**



19

Memory Locking: Feature summary

- Attestation process can be interrupted at any time (@block granularity)
- Memory gradually locked or unlocked during attestation
 - Guarantees temporal consistency
 - Detects transient and self-relocating malware
- Requires hardware support to lock/unlock memory blocks
 - MMU (as in HYDRA)
 - Runtime-configurable MPU (as in TyTan)
- Other approaches possible, e.g., inconsistency detection via dirty bits

20

Shuffled Measurements [HOST'18]

- Allows attestation process to be interrupted
- Memory traversed in unpredictable AND secret fashion

21

Shuffled Measurements

1	2	3	4	5	6
---	---	---	---	---	---

Permutation = {3, 2, 5, 6, 1, 4}

Shuffled Measurements

1	2	3	4	5	6
---	---	---	---	---	---

Permutation = {3, 2, 5, 6, 1, 4}

Shuffled Measurements

1	2	3	4	5	6
---	---	---	---	---	---

Permutation = {3, 2, 5, 6, 1, 4}

Shuffled Measurements




Permutation = {3, 2, 5, 6, 1, 4}

Shuffled Measurements



Permutation = {3, 2, 5, 6, 1, 4}

Shuffled Measurements



Permutation = {3, 2, 5, 6, 1, 4}

Shuffled Measurements



Permutation = {3, 2, 5, 6, 1, 4}

Shuffled Measurements [HOST'18]

- Allow attestation to be interrupted @block granularity
- Memory traversed in *unpredictable AND secret* fashion
- One-time traversal permutation generated by PRF on input of:
Key, counter and/or clock, challenge
- Permutation protected by:
 - Storing it in separate (secure) working memory → new feature
OR
 - Storing encrypt-then-mac of each permutation element in regular memory
- For large number of blocks, **37% malware evasion** rate (see paper)
 - Multiple independent measurements to reliably detect malware

29

Shuffled Measurements: Summary

- Attestation process can be interrupted
- Measures memory in random & secret order
 - Probabilistic detection of self-relocating malware
- Multiple measurements to ensure high detection rate
 - Extra overhead in attestation run-time
- Needs no additional hardware features (over basic RA)
 - Encrypt-then-mac each element in permutation

30

Self-Measurement


- **Prover** periodically measures own memory:
 - Current RA techniques work “on-demand”, i.e., based on **Verifier**’s explicit request

BUT:

- Prover knows best → can schedule measurements to avoid interference with safety-critical tasks
- Attestation process remains uninterrupted
- Two recent techniques:
 - SeED [WiSec’17]
 - ERASMUS [HOST’18]

31

Self-Measurements: On-demand RA




32

Self-Measurements: SeED [WiSec'17]

- Idea: Prover measures itself and immediately sends result to Verifier without any explicit request by the latter
- Unpredictable measurement schedule to protect against transient malware
 - Requires secure real-time clock

33

Self-Measurements: SeED [WiSec'17]



34

Self-Measurements: ERASMUS [DATE'18]

- Similar to SeED, but:

Instead of sending one result to Verifier after each self-measurement:


- Accumulates measurements locally
- Sends **all** (or specified range) to Verifier upon request

- Still requires secure real-time clock to detect transient malware


35

Self-Measurements: ERASMUS [DATE'18]

Measurement Phase



Collection Phase



36

Self-measurements: Feature summary

- No presence of Verifier
 - Suitable for unattended settings
- Unpredictable RA schedule
 - Detects transient malware
- Uninterruptible attestation
 - Detects self-relocating malware
- Additional hardware feature
 - Secure (i.e., reliable, read-only) real-time clock

37

Conclusions

- Tension between secure RA and safety-critical operation
- Surveyed RA techniques attempt to mitigate it by:
 - Memory locking + temporal consistency
 - Shuffled memory traversal
 - Self-measurements
- Identified trade-offs between:
 - Malware detection
 - Temporal consistency guarantees
 - Run-time overhead
 - Hardware requirements

38

Ongoing & Future Work

- Formal verification
- Swarm/group setting
- Extensions (actually, practical applications), e.g.:
 - Secure Reset, Erasure, SW/FW Update

39

\end{document}

% Comments?
% Questions?

40

MITIGATION METHODS	Malware Types Detected		Writable Memory Availability	Consistency Guarantees	Interruptibility	Unattended Setting	Extra HW Require	Run-Time Overhead
	Self-relocating	Transient						
TyTan [DAC'15]			Red	Green	Red	Red	Red	Green
Temporal Consistency via Memory Locking [ASIACCS'18]			Green	Green	Green	Red	Red	Green
Shuffled Measurements SMARM [HOST'18]	Green	Red	Green	Red	Green	Red	Green	Red
Self-measurements SeED, ERASMUS [WiSec'17, DATE'18]	Green	Green	Red	Green	Red	Green	Red	Green

41